Student Performance Analysis Using Machine Learning Techniques

Nasser Hasan Padilla

University of the Cumberlands

Advanced Big Data and Data Mining (MSCS-634-B01)

Professor Satish Penmatsa

April 24, 2026

Abstract

This project explores student performance data using machine learning techniques to identify patterns and predict academic outcomes. The study applies data preprocessing, exploratory data analysis (EDA), regression modeling, classification, clustering, and association rule mining. The results demonstrate that early academic indicators and behavioral factors significantly influence student performance. The findings provide insights that can be used to identify at-risk students and improve educational outcomes. Ethical considerations such as bias and data privacy are also discussed.

Introduction

The purpose of this project is to analyze student academic performance using data mining and machine learning techniques. The dataset used in this study is the Student Performance dataset obtained from the UCI Machine Learning Repository. This dataset contains a variety of features related to student demographics, study habits, and academic results.

The objective of this analysis is to uncover relationships within the data and build predictive models that can estimate student performance. By leveraging multiple analytical techniques, including regression, classification, clustering, and association rule mining, this study aims to provide actionable insights that can be applied in real-world educational settings.

Dataset Description

The dataset includes student information such as age, gender, family background, study time, and prior academic performance. The primary target variable for this study is the final grade (G3), which represents the student's overall academic outcome.

This dataset was selected because it provides a comprehensive view of factors that influence academic success and supports multiple types of machine learning analysis.

Data Preparation and Preprocessing

Data preprocessing was performed to ensure the dataset was suitable for analysis. The following steps were taken:

- Categorical variables were encoded using label encoding techniques.
- The dataset was inspected for missing values, and none were found that required removal or imputation.
- Numeric features were standardized where necessary to improve model performance.
- Feature selection was performed to focus on variables most relevant to predicting student performance.

These steps ensured that the dataset was clean, consistent, and ready for modeling.

Exploratory Data Analysis (EDA)

Exploratory Data Analysis was conducted to identify patterns and relationships within the dataset.

Key findings include:

- A strong positive correlation between early grades (G1 and G2) and final grade (G3).
- Study time shows a positive relationship with academic performance.
- Absences and alcohol consumption exhibit a negative correlation with final grades.

Figure 1

Correlation Heatmap of Features

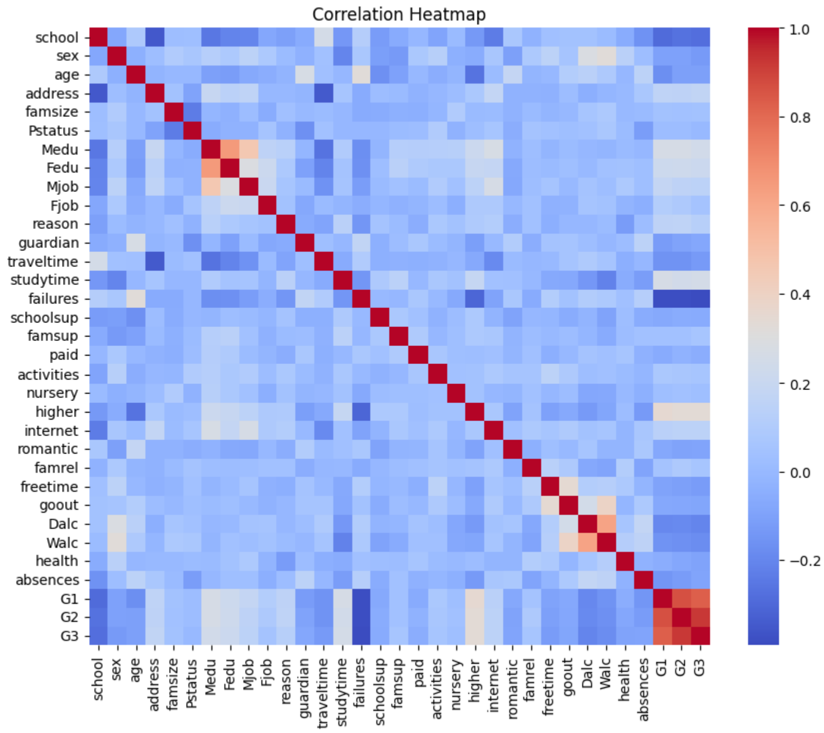

Figure 2

Distribution of Final Grades (G3)

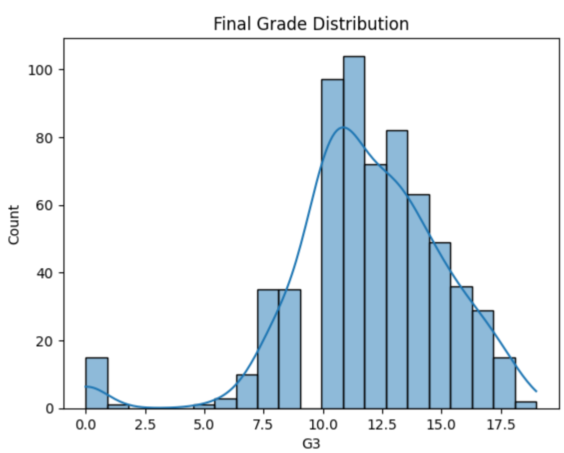

These insights guided the feature selection process and informed the modeling approach.

Feature Engineering

Feature engineering was applied to improve model performance and interpretability. A classification target variable (G3_class) was created by grouping final grades into categories representing low, medium, and high performance levels.

This transformation enabled the use of classification models and provided a clearer understanding of student performance segments.

Regression Modeling

Regression models were developed to predict the final grade (G3).

Models Used:

- Linear Regression
- Ridge Regression

Evaluation Metrics:

- R-squared (R^2)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)

Results:

Ridge Regression demonstrated slightly improved performance compared to Linear Regression due to its ability to handle multicollinearity.

Figure 3

Regression Model Performance Comparison

```
R2: 0.8202961916526714
MSE: 1.3093921278041443
```

Classification Modeling

Classification models were developed to categorize student performance levels.

Models Used:

- Decision Tree Classifier
- k-Nearest Neighbors (k-NN)

Evaluation Metrics:

- Accuracy Score
- Confusion Matrix
- F1 Score

The k-NN model achieved higher accuracy after tuning, while the Decision Tree provided better interpretability.

Figure 4

Confusion Matrix for Classification Model

```
Accuracy: 0.8307692307692308

Confusion Matrix:
 [[23  9  0]
 [ 5 66  5]
 [ 0  3 19]]

Classification Report:
              precision    recall  f1-score   support

           0       0.82      0.72      0.77        32
           1       0.85      0.87      0.86        76
           2       0.79      0.86      0.83        22

    accuracy                           0.83       130
   macro avg       0.82      0.82      0.82       130
weighted avg       0.83      0.83      0.83       130
```

Clustering Analysis

K-Means clustering was applied to group students based on similar characteristics.

Findings:

- Students were grouped into clusters representing different performance levels.
- Clear distinctions were observed between high-performing and low-performing students.

Figure 5

Student Clusters Visualization

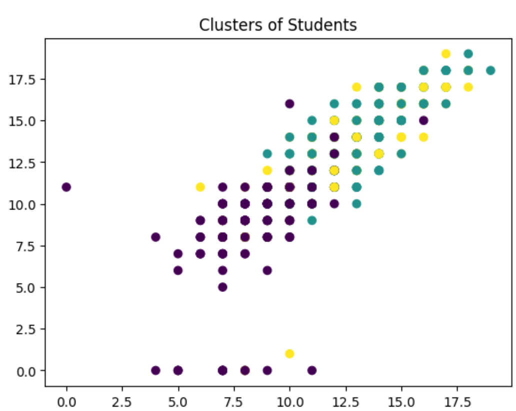

This analysis helps identify groups of students that may require targeted interventions.

Association Rule Mining

Association rule mining was performed using the Apriori algorithm to identify patterns in student behavior.

Key Insights:

- Students with low study time and high absences are more likely to perform poorly.
- Positive academic habits are frequently associated with higher grades.

Practical Recommendations

Based on the analysis, the following recommendations are proposed:

1. Monitor early academic performance to identify at-risk students.
2. Encourage consistent study habits to improve outcomes.
3. Reduce absenteeism through targeted interventions.
4. Use predictive models to support academic advising systems.

These recommendations can help educational institutions improve student success rates.

Ethical Considerations

Several ethical considerations were addressed in this project:

- **Data Privacy:** Student data must be handled securely and anonymized to protect personal information.

- **Bias:** The dataset may reflect socioeconomic or demographic biases, which can affect model predictions.
- **Fairness:** Predictive models should not be used to unfairly label or disadvantage students.

Efforts were made to interpret results responsibly and avoid biased conclusions.

Conclusion

This project demonstrates the effectiveness of machine learning techniques in analyzing student performance data. By combining regression, classification, clustering, and association rule mining, meaningful insights were obtained.

The findings highlight the importance of early academic indicators and behavioral factors in predicting student success. Future work could involve using larger datasets, incorporating additional features, and exploring advanced machine learning models.

References

Cortez, P., & Silva, A. (2008). Using data mining to predict secondary school student performance. UCI Machine Learning Repository.